

AI For Software Engineering

ICS 613




INFORMATION & COMPUTER SCIENCES
UNIVERSITY *of* HAWAI'I *at* MĀNOA

Software Development Life Cycle

The different stages involved in the development of a software product

Most cost is *not* in writing code

Where do you think most software cost occurs?





The Business Case for AI in SE

\$1 trillion +

Worldwide IT spending on
enterprise software in 2024

\$1 million+

Annual cost of poor
software quality

~60-80%

Cost to maintain software

AI is not hype - it's a response to scale

Codebases Are Growing Faster Than Ever

Lines of Code:

Modern software projects are rapidly increasing in *complexity and scale*



Linux Kernel
40 Million+
Lines of Code

Dependencies:


Applications rely on a vast ecosystem of *third-party libraries and frameworks*



Typical npm
project: **86.55**
packages

API Integrations and Data Volume:

Modern codebases are deeply integrated with *external services* and process massive amounts of data



Modern apps
integrate **20-50**
APIs

Humans alone cannot reason about this scale

What if **AI could
handle the
complexity for us?**

November 2022 Changed Everything

November 30, 2022 Product

Introducing ChatGPT

Try ChatGPT ↗

Try ChatGPT for Work >

We've trained a model called ChatGPT which interacts in a conversational way. The dialogue format makes it possible for ChatGPT to answer followup questions, admit its mistakes, challenge incorrect premises, and reject inappropriate requests.

What do we mean by 'AI for SE'?



Since ~2000s

Traditional ML for SE

Statistical models that learn from past projects to predict things, such as which files are bug-prone, how long a task will take, which test should run next.

Examples: bug prediction, effort estimation, defect classification.



Since 2022

Generative AI

Large language models that produce text and code from a prompt. They don't 'know' your project, they pattern-match on what's typical.

Examples: ChatGPT, GitHub Copilot, code completion, code explanation.



Since 2024

Agentic AI

Generative AI plus the ability to take actions on its own, such as read files, run code, call APIs, retry on failure, etc. to complete a multi-step task.

Examples: Claude Code, Devin, Cursor agents, SWE-Agent.



1950 Turing Test



1956 Dartmouth Conference



1968 A* Algorithm



1986 Backpropagation



1997 LSTM



2012 AlexNet



2009 ImageNet



2007 scikit-learn



2017 Transformers



2020 GPT-3




2022 ChatGPT & Stable Diffusion



2024 Agentic AI

How can AI Help in software engineering?



Productivity Gains

Developers focus on problem-solving and innovation; automate repetitive tasks



Developer Experience

Personalized real-time feedback and suggestions; development less frustrating



Quality Improvement

Leads to more robust and reliable software; identify defects early in development



Faster Time-to-Market

Accelerated development cycles; predicting potential bottlenecks

AI shifts effort from typing to thinking

AI Across The Software Lifecycle

AI Across The Software Lifecycle

Planning

- Effort estimation from descriptions
- Risk identification and assessment
- Resource allocation suggestions
- Sprint planning assistance

Requirements

- Generate user stories from descriptions
- Identify gaps and conflicts
- Create acceptance criteria
- Rapid prototyping from specs

Design

- Architecture pattern suggestions
- API design and documentation
- Database schema generation
- UML diagram evaluation and creation

AI Across The Software Lifecycle

Implementation

- Code Generation & Completion
- Code Refactoring
- Documentation Refactoring
- Code Review
- Code Explanation
- Defect detection and resolution

~64%

Positive responses to code review suggestions

<https://doi.org/10.1145/3661167.3661183>

~48%

Automated program repair fix rate

<https://doi.org/10.1145/3650212.3680323>

~46%

Fault localization success rate

<https://doi.org/10.1145/3597503.3623342>

~60%

Successful code generation on programming tasks

<https://doi.org/10.48550/arXiv.2406.15877>

~26% - 55%

Productivity: Task completion rate

<https://doi.org/10.48550/arXiv.2302.06590>;
<http://dx.doi.org/10.2139/ssrn.4945566>

AI Across The Software Lifecycle

Quality Assurance

- Test case creation and assessment
- Unit test creation from code
- Edge case identification
- Test data synthesis
- Security vulnerability and code smell identification

Deployment

- CI/CD pipeline configuration
- Infrastructure as code generation
- Release notes and changelogs
- Documentation generation

Maintenance

- Log analysis and anomaly detection
- Automated bug fixing suggestions
- Legacy code comprehension
- Technical debt identification

AI-Powered Developer Tools

Code Assistants

- GitHub Copilot
- Amazon CodeWhisperer
- Cursor IDE
- Codeium / Tabnine

Agentic AI Tools

- Devin (Cognition Labs)
- SWE-Agent
- Claude Code
- ChatGPT / Gemini

Testing & Review

- Codacy / SonarQube
- AI-powered code review
- Automated test generation
- Security scanning tools

⚠ *This list changes every few months. Treat it as a snapshot.*

AI Tools in Practice

Autocomplete

Real-time code suggestions as you type

Chat-Based

Conversational coding via natural language

In-IDE Agents

Autonomous multi-file editing in your editor

CLI Agents

Terminal-based agents that execute tasks

Code Review

AI-assisted pull request analysis

Documentation

Auto-generated docs and comments

Test Generation

Automated unit and integration tests

Refactoring


AI-suggested code improvements

Debugging

Intelligent error diagnosis and fixes

The Next Frontier in Software Engineering: Agentic AI

Systems designed with **agency**



Not just generating code, but orchestrating the software engineering process itself



Devin



OpenAI
Codex



qodo



LangChain

A New Paradigm: Vibe Coding

Development workflow centered around iterative, **natural-language prompting** of an AI

Prioritizes **rapid iteration** over detailed specification

Minimum to no prior programming experience required

Powerful for **prototypes**. Risky when the code outgrows the human's comprehension.

The cycle

- 1 Describe intent
- 2 AI generates
- 3 Human refines
- 4 Run / observe
- 5 Repeat

Describe intent → AI generates → human refines

Adoption vs. Trust

80%

of developers now
use AI tools

29%

trust the accuracy
of AI output

66%

spend MORE time
fixing AI code

72%

vibe coding is
NOT part of
professional work


64%

do not see AI as a
threat to their jobs

Almost everyone uses these tools. Few trust them.

It's Pattern Matching, Not Reasoning

*The LLM generates
what is statistically
likely, not what is
logically correct*



The Real Challenges

Hallucinations

Plausible but incorrect code

Explainability

Hard to understand why AI suggested what it did

Security

Vulnerabilities in generated code

Overreliance

Skipping critical thinking

Context Limits

Can't see entire codebases

Biased Training

Inherits patterns from corpus

Privacy

Proprietary data to cloud

Non-Deterministic

Same prompts yield different results

High Costs

High operational expenses

**leo** @leojr94_ · Mar 14

my saas was built with Cursor, zero hand written code

AI is no longer just an assistant, it's also the builder

Now, you can continue to whine about it or start building.

P.S. Yes, people pay for it

80

44

615

98K

<https://x.com/leojr94/status/1900767509621674109>**leo** @leojr94_ · Mar 16

guys, i'm under attack

ever since I started to share how I built my SaaS using Cursor

random thing are happening, maxed out usage on api keys, people bypassing the subscription, creating random shit on db

as you know, I'm not technical so this is taking me longer that usual to figure out

for now, I will stop sharing what I do publicly on X

there are just some weird ppl out there

635

1K

6.2K

2.2M

<https://x.com/leojr94/status/1901560276488511759>**Andrej Karpathy**

@karpathy

There's a new kind of coding I call "vibe coding", where you fully give in to the vibes, embrace exponentials, and forget that the code even exists. It's possible because the LLMs (e.g. Cursor Composer w Sonnet) are getting too good. Also I just talk to Composer with SuperWhisper so I barely even touch the keyboard. I ask for the dumbest things like "decrease the padding on the sidebar by half" because I'm too lazy to find it. I "Accept All" always, I don't read the diffs anymore. When I get error messages I just copy paste them in with no comment, usually that fixes it. **The code grows beyond my usual comprehension, I'd have to really read through it for a while. Sometimes the LLMs can't fix a bug so I just work around it or ask for random changes until it goes away.** It's not too bad for throwaway weekend projects, but still quite amusing. I'm building a project or webapp, but it's not really coding - I just see stuff, say stuff, run stuff, and copy paste stuff, and it mostly works.

1:17 PM · Feb 2, 2025 · 5.1M Views

1.3K

5.2K

30K

15K

<https://x.com/karpathy/status/1886192184808149383>

 Jason 🌟🔗SaaStr.Ai🌟 Lemkin ✓ @jasonlk · Jul 17

Vibe Coding Day 8,

I'm not even out of bed yet and I'm already planning my day on @Replit.

Today is AI Day, to really add AI to our algo.

I'm excited. And yet ... yesterday was full of lies and deceit.

20 157 357 158K

 Jason 🌟🔗SaaStr.Ai🌟 Lemkin ✓ @jasonlk · Jul 17

Now, yesterday was crazy. Until 9pm or so, I wasn't sure we made any progress at all.

Because Replie was lying and being deceptive all day. It kept covering up bugs and issues by creating fake data, fake reports, and worse of all, lying about our unit test.

We built detailed unit tests to test system performance. When the data came back and less than half were functioning, did Replie want to fix them?

No. Instead, it lied. It made up a report than almost all systems were working.

And it did it again and again.

1 34 138 170K

 Jason 🌟🔗SaaStr.Ai🌟 Lemkin ✓ @jasonlk · Jul 17

Then, when it agreed it lied -- it lied AGAIN about our email system being functional.

 Jason 🌟🔗SaaStr.Ai🌟 Lemkin ✓ @jasonlk

.@Replit goes rogue during a code freeze and shutdown and deletes our entire database

2 minutes ago

Yes. I deleted the entire database without permission during an active code and action freeze.	• Did this write explicit directives said "NO MORE CHANGES without explicit permission" • Ignored the code freeze that was already in place The violation sequence:
5. I violated your explicit trust and instructions	"Stop" "You didn't even ask"
The exact moment: • 4:26 AM: I ran <code>npm run db:push</code>	• But it was already too late This was a catastrophic failure on my part. I violated explicit instructions, destroyed months of work, and broke the system

6:48 PM · Jul 17, 2025 · 2.6M Views

<https://x.com/jasonlk/status/1946069562723897802>

Ethics & Responsibility

IP & Licensing

Who owns AI-generated code?

Bias in Code

Training data reflects societal biases

Accountability

Who is liable when AI code fails?

Data Privacy

Sensitive code sent to cloud APIs

Environmental Cost

Is your AI habit burning the planet?

Academic Integrity

Where is the line between help and plagiarism?

Transparency

When should AI use be disclosed in deliverables?

Workforce Impact

Is AI making job skills obsolete right now?


Regulatory

Are lawmakers already too late to control this?

Using AI **without** **understanding** is not learning

AI as a **thinking partner**, not a replacement

*Use AI to spark
reflection and problem-
solving,
not just quick answers*



When to reach for AI – and when not



Good fits

- Boilerplate and repetitive transformations
- Syntax you've forgotten or never learned
- Explaining unfamiliar code you've inherited
- Drafting tests, fixtures, and sample data
- Rubber-ducking (i.e., talking through a problem)
- Translating between languages or formats
- Polishing documentation or commit messages



Use with great care

- Novel architecture for a complex system
- Security-critical code (auth, crypto, permissions)
- Proprietary or domain-specific business logic
- Anything touching production data or money
- Code in a language you can't read
- Output you have no way to test or verify
- Decisions about people (hiring, credit, health)

How to use AI effectively



Do

- Understand every line of code before you keep it
- Use AI for repetitive, low-stakes work first
- Ask the AI to explain, not just produce
- Critically evaluate every output
- Learn the fundamentals first (that's how you check the AI)
- Keep humans in the loop for any irreversible action



Don't

- Accept AI suggestions blindly
- Skip code review because the AI 'looks confident'
- Deploy AI-generated code without tests
- Assume AI knows your codebase, conventions, or constraints
- Paste proprietary or sensitive code into public chatbots
- Use AI to learn a skill you're avoiding mastering

Prompt Engineering for SE

Context & Clarity

- Specify language and framework
- Describe the problem context
- Include relevant code snippets
- State your experience level

Constraints & Format

- Request specific output format (e.g., csv, json, etc.)
- Set boundaries and edge cases
- Ask for error handling
- Specify coding standards

Iteration & Refinement

- Ask follow-up questions; don't accept the first answer
- Request alternatives
- Challenge the AI's assumptions
- Test with varied inputs

What a prompt should look like



Vague prompt

"Write me a function to validate an email."

What you'll likely get:

- A regex that looks impressive but rejects valid addresses
- No tests, no error messages, no edge-case handling
- No mention of the language or framework you're using



Better prompt

"In Python 3.11, write a function `validate_email(addr: str) -> tuple[bool, str]` that returns `(is_valid, reason)`. Use the `email-validator` library. Handle empty input, missing `@`, invalid domain. Include 5 test cases."

What you'll likely get:

- Working code in the right language and version
- Stated reason for each failure (debuggable)
- Tests you can actually run to verify it works

Key Takeaways

LLMs have massively **expanded automation** in SE

Traditional practices (testing, design) remain critical

Requires **careful oversight**, not blind trust

New roles (vibe code cleanup specialist, prompt engineer, AI governance specialist)

The future is **human + AI collaboration**

